

# Recurrent Neural Network

In this task, we implement RNN cells to understand the computation of RNN. Then we build RNN with different cells for a language modeling task.

```
In [1]: # As usual, a bit of setup
import time
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-i
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

## Recurrent Neural Networks

### A toy problem

```
In [2]: ## Setup an example. Provide sizes and the input data.

# set sizes
time_steps = 5
batch_size = 4
input_size = 3
hidden_size = 2

# create input data with shape [batch_size, time_steps, num_features]
np.random.seed(15009)
input_data = np.random.rand(batch_size, time_steps, input_size).astype(np.f
```

### Implement an RNN and a GRU with tensorflow

```

In [3]: ## Create an RNN model
tf.reset_default_graph()
tf.random.set_random_seed(15009)

# initialize a state of zero for both RNN and GRU
# 'state' is a tensor of shape [batch_size, hidden_size]
init_state = np.zeros([batch_size, hidden_size])
initial_state = tf.Variable(init_state, dtype=tf.float32)

# create a BasicRNNCell
rnn_cell = tf.nn.rnn_cell.BasicRNNCell(hidden_size)

# 'outputs' is a tensor of shape [batch_size, max_time, hidden_size]
# RNN cell outputs the hidden state directly, so the output at each step is
# final_state is the last state of the sequence. final_state == outputs[:,
rnn_outputs, rnn_final_state = tf.nn.dynamic_rnn(rnn_cell, input_data,
                                                  initial_state=initial_state,
                                                  dtype=tf.float32)

# create a GRUCell
gru_cell = tf.nn.rnn_cell.GRUCell(hidden_size)

# 'outputs' is a tensor of shape [batch_size, time_steps, hidden_size]
# Same as the basic RNN cell, final_state == outputs[:, -1, :]
gru_outputs, gru_final_state = tf.nn.dynamic_rnn(gru_cell, input_data,
                                                  initial_state=initial_state,
                                                  dtype=tf.float32)

# initialize variables
init = tf.global_variables_initializer()
session = tf.Session()
session.run(init)

# run the RNN model and get outputs and the final state
tfrnn_outputs, tfrnn_final_state = session.run([rnn_outputs, rnn_final_state])

# run the GRU model and get outputs and the final state
tfgru_outputs, tfgru_final_state = session.run([gru_outputs, gru_final_state])

```

WARNING:tensorflow:From /Users/thomasklimek/anaconda3/lib/python3.7/site-packages/tensorflow/python/framework/op\_def\_library.py:263: colocate\_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From <ipython-input-3-40b6cc66db36>:11: BasicRNNCell.\_\_init\_\_ (from tensorflow.python.ops.rnn\_cell\_impl) is deprecated and will be removed in a future version.

Instructions for updating:

This class is equivalent as tf.keras.layers.SimpleRNNCell, and will be replaced by that in Tensorflow 2.0.

WARNING:tensorflow:From <ipython-input-3-40b6cc66db36>:19: dynamic\_rnn (from tensorflow.python.ops.rnn) is deprecated and will be removed in a future version.

ure version.

Instructions for updating:

Please use `keras.layers.RNN(cell)`, which is equivalent to this API

WARNING:tensorflow:From <ipython-input-3-40b6cc66db36>:23: GRUCell.\_\_init\_\_ (from tensorflow.python.ops.rnn\_cell\_impl) is deprecated and will be removed in a future version.

Instructions for updating:

This class is equivalent as tf.keras.layers.GRUCell, and will be replaced by that in Tensorflow 2.0.

## Read out parameters from RNN and GRU cells

```
In [4]: from rnn_param_helper import get_rnn_params, get_gru_params

#####
# 0. Understand TF BasicRNN and GRU parameters (0 points)
# Please read the code and documentation of get_rnn_params and get_gru_params
# what are these parameters. You will need to use these parameters in your
# NO implementation is needed here.
#####

wt_h, wt_x, bias = get_rnn_params(rnn_cell, session)
wtu_h, wtu_x, biasu, wtr_h, wtr_x, biasr, wtc_h, wtc_x, biasc = get_gru_params(rnn_cell, session)
```

## Numpy Implementation

Implement your own RNN model with numpy. Your implementation needs to match the tensorflow calculation.

```
In [6]: #####
# 1. Implementation (9 points)
# Implement your basic RNN and GRU RNN in implementations/rnn.py and compare
# tensorflow functions
#####

from implementations.rnn import rnn,gru

# calculation from your own implemenation of a basic RNN
nprnn_outputs, nprnn_final_state = rnn(wt_h, wt_x, bias, init_state, input_

print("Difference between your RNN implementation and tf RNN",
      rel_error(tfrnn_outputs, nprnn_outputs) + rel_error(tf

# calculation from your own implemenation of a GRU RNN
npgru_outputs, npgru_final_state = gru(wtu_h, wtu_x, biasu, wtr_h, wtr_x, b

print("Difference between your GRU implementation and tf GRU",
      rel_error(tfgru_outputs, npgru_outputs) + rel_error(tfgru_final_state
```

Difference between your RNN implementation and tf RNN 2.1863736249496377e-07

Difference between your GRU implementation and tf GRU 0.9037854511656668

## GRU includes RNN as a special case

Can you assign a special set of parameters to GRU such that its outputs is almost the same as RNN?

```
In [7]: # Assign some value to a parameter of GRU

#####
# 2. Setting GRU weights (4 points)
# Get weights/bias from the basic RNN and set them to some GRU weights/bias
# Then set some other parameter of GRU, then GRU recovers RNN.
#####

#session.run(gru_cell.weights[0].assign(?))
#session.run(gru_cell.weights[1].assign(?))
#session.run(gru_cell.weights[2].assign(?))
#session.run(gru_cell.weights[3].assign(?))

# outputs from the GRU with special parameters.
updated_outputs = session.run(gru_outputs)

# they are the same as the calculation from the basic RNN
print("Difference between RNN and a special GRU", rel_error(tfrnn_outputs,
```

Difference between RNN and a special GRU 1.0

## Long term dependency: forward

In this experiment, you will see that the basic RNN model is hard to keep long term dependency

```

In [346]: from rnn_param_helper import set_rnn_params, set_gru_params

# Create a larger problem

# set sizes
time_steps = 50
batch_size = 100
input_size = 5
hidden_size = 8

# create input data with shape [batch_size, time_steps, num_features]
np.random.seed(15009)
input_data = np.random.rand(batch_size, time_steps, input_size).astype(np.f

## Create an RNN model with GRU
tf.reset_default_graph()
tf.random.set_random_seed(15009)

# copy the basic RNN and the GRU RNN above here:

initial_state = tf.Variable(np.zeros([batch_size, hidden_size]), dtype=tf.f

#####
# 3. Apply TF RNN functions (2 points)
# Please use the tensorflow function for the basic RNN and the GRU RNN belo
# from the larger problem. Basically you just need to copy some code above
#####

#rnn_outputs, _ = tf.nn.dynamic_rnn(rnn_cell, ...) # please complete this l
#gru_outputs, _ = tf.nn.dynamic_rnn(gru_cell, ...) # please complete this l

# initialize variables
init = tf.global_variables_initializer()
session = tf.Session()
session.run(init)

def show_hist_of_hidden_values(session, initial_state, state, title):
    """Set `initial_state` to different values and run the `state` value. C
    values due to different initializations. If the model cannot capture
    initialization does not have much effect to the value of `state` at
    """

    batch_size, hiddens_size = state.get_shape()

    # intialize the model with different initial states and then calculate
    init_zero = np.zeros([batch_size, hidden_size])
    session.run(initial_state.assign(init_zero))
    state_zero_init = session.run(state)

    init_rand1 = np.random.rand(batch_size, hidden_size)
    session.run(initial_state.assign(init_rand1))
    state_rand1_init = session.run(state)

    init_rand2 = np.random.rand(batch_size, hidden_size)

```

```

session.run(initial_state.assign(init_rand2))
state_rand2_init = session.run(state)

init_scaleup1 = init_rand1 * 100
session.run(initial_state.assign(init_scaleup1))
state_scaleup1_init = session.run(state)

# plot the difference between the four difference settings

# For each sequence, calculate the norm of the difference of the states
norm_diff1 = np.linalg.norm(state_zero_init - state_rand1_init, axis=1)
norm_diff2 = np.linalg.norm(state_zero_init - state_rand2_init, axis=1)
norm_diff3 = np.linalg.norm(state_rand1_init - state_rand2_init, axis=1)
norm_diff4 = np.linalg.norm(state_scaleup1_init - state_zero_init, axis=1)
norm_diff5 = np.linalg.norm(state_scaleup1_init - state_rand1_init, axis=1)
norm_diff6 = np.linalg.norm(state_scaleup1_init - state_rand2_init, axis=1)

# plot the histogram of norms of differences
n_bins = 20
fig, axs = plt.subplots(2, 3, sharey=True, tight_layout=True)
plt.suptitle(title, fontsize=16)

axs[0, 0].hist(norm_diff1, bins=n_bins)
axs[0, 1].hist(norm_diff2, bins=n_bins)
axs[0, 2].hist(norm_diff3, bins=n_bins)
axs[1, 0].hist(norm_diff4, bins=n_bins)
axs[1, 1].hist(norm_diff5, bins=n_bins)
axs[1, 2].hist(norm_diff6, bins=n_bins)

# set values for the basic RNN model

# play with the scale, and see if you can find any value that achieves long
scale = 2.0
wt_h = (np.random.rand(hidden_size, hidden_size) - 0.5) * scale
wt_x = (np.random.rand(input_size, hidden_size) - 0.5) * scale
bias = (np.random.rand(hidden_size) - 0.5) * scale

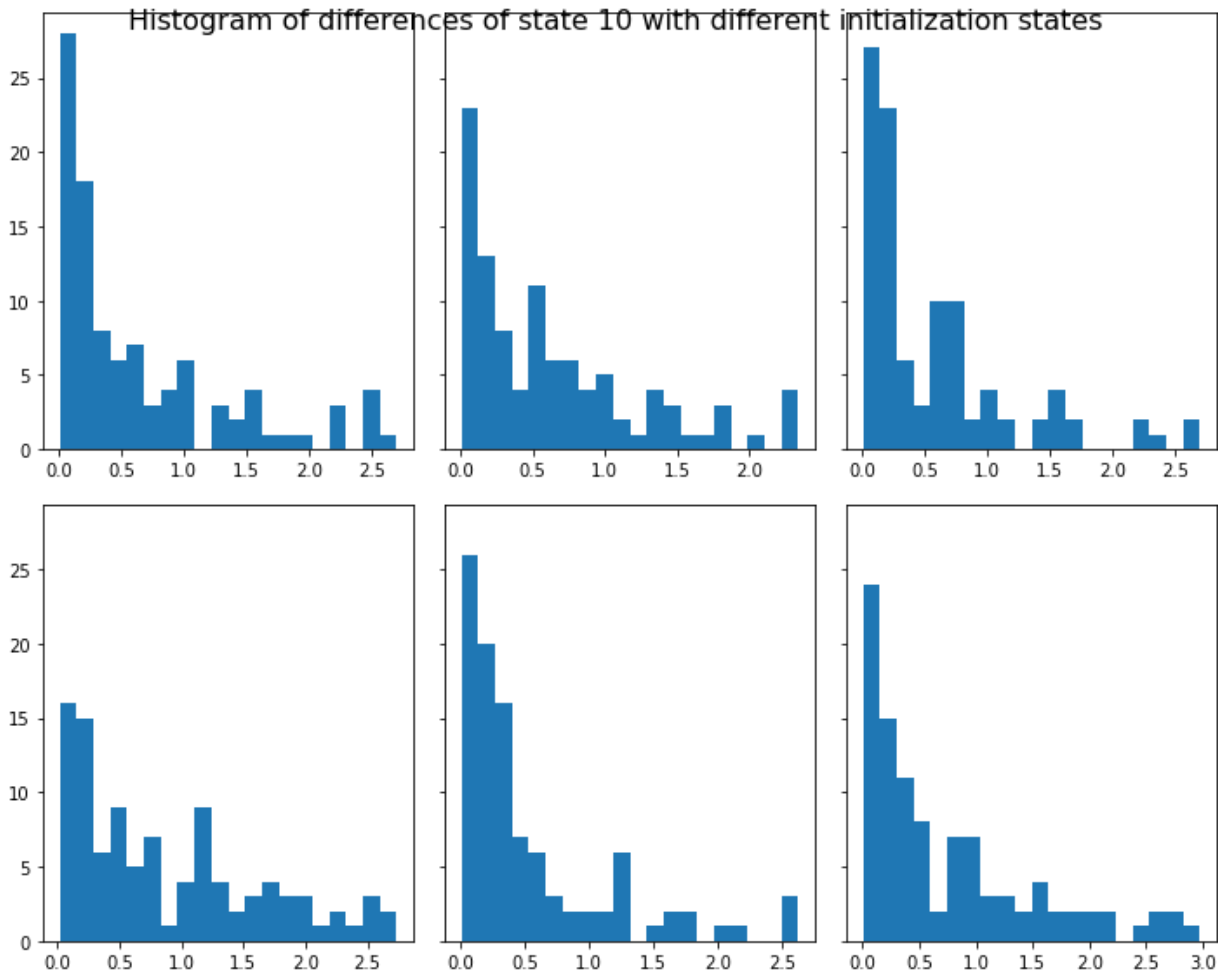
set_rnn_params(rnn_cell, session, wt_h, wt_x, bias)

# get the 10th state and check its dependency on the initial state
rnn_state10 = tf.transpose(rnn_outputs, [1, 0, 2])[10]
show_hist_of_hidden_values(session, initial_state, rnn_state10,
                           'Histogram of differences of state 10 with diffe

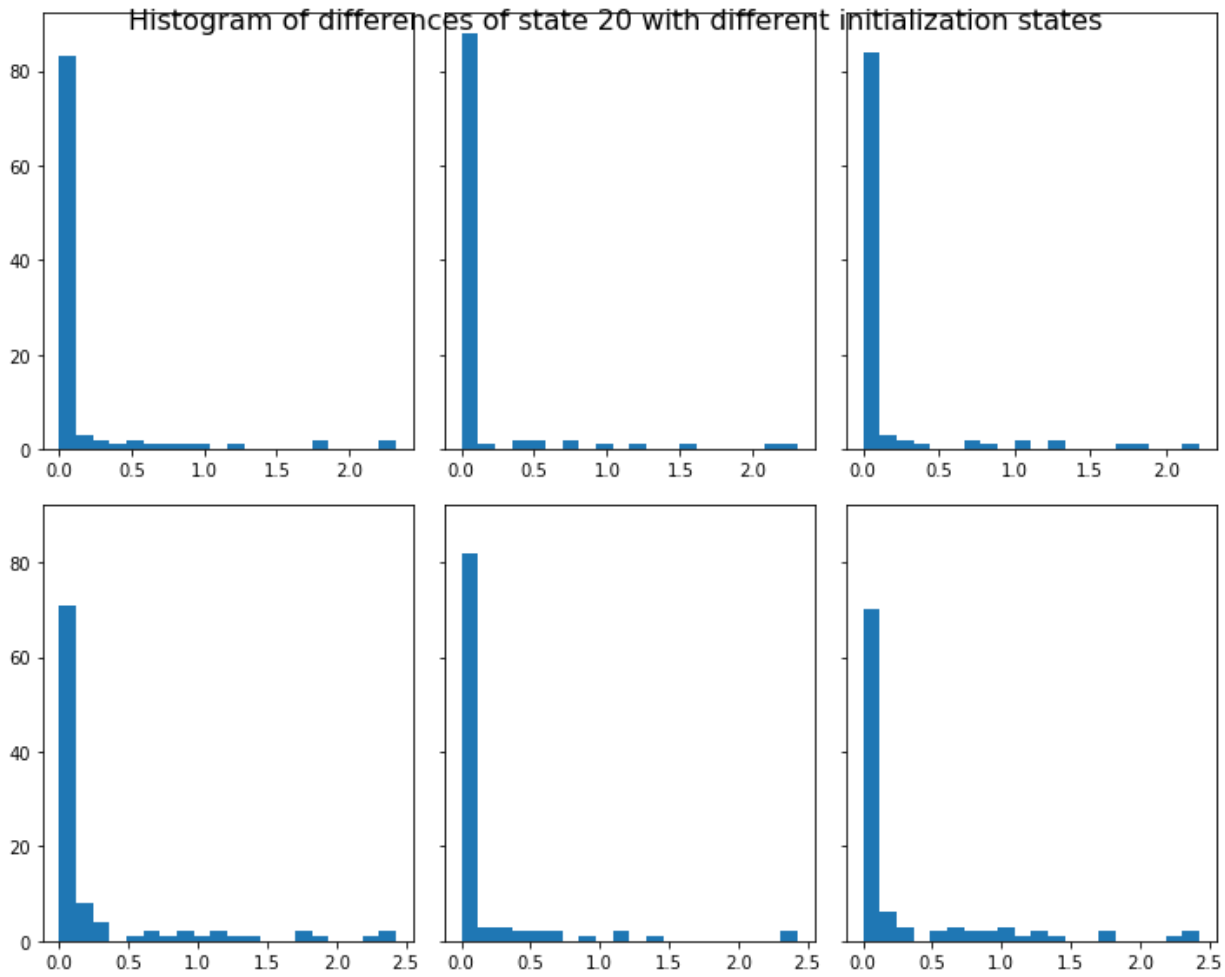
# get the 20th state and check its dependency on the initial state
rnn_state20 = tf.transpose(rnn_outputs, [1, 0, 2])[20]
show_hist_of_hidden_values(session, initial_state, rnn_state20,
                           'Histogram of differences of state 20 with diffe

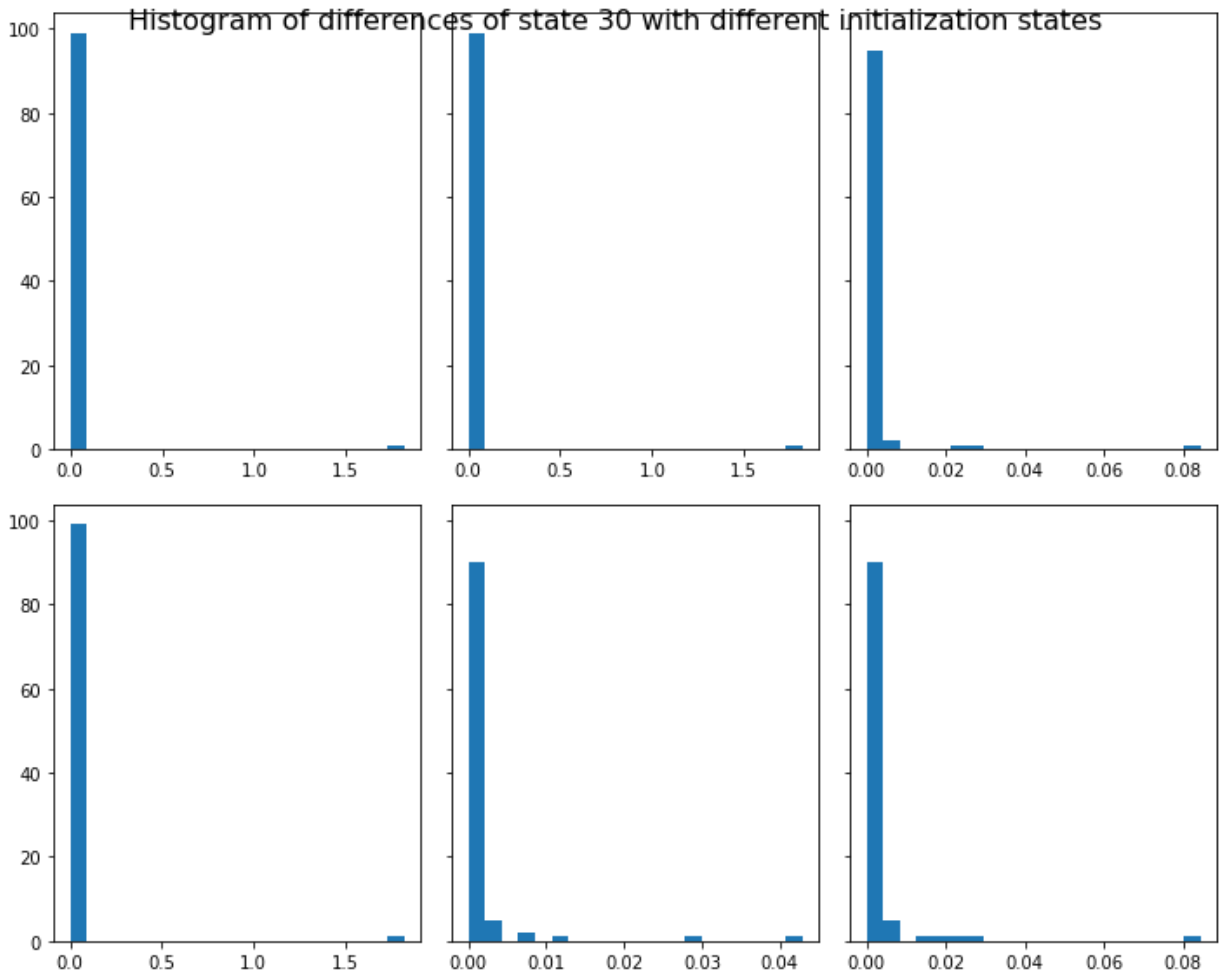
# get the 30th state and check its dependency on the initial state
rnn_state30 = tf.transpose(rnn_outputs, [1, 0, 2])[30]
show_hist_of_hidden_values(session, initial_state, rnn_state30,
                           'Histogram of differences of state 30 with diffe

```









```

In [347]: # Can you set GRU parameters such that it maintains the initial state in th

scale_gru = 5.0


#####
# 4. Setting GRU parameters (4 points) #
# Set GRU parameters here so that it can capture long term dependency #
#####


# get the 10th state
gru_state10 = tf.transpose(gru_outputs, [1, 0, 2])[10]
show_hist_of_hidden_values(session, gru_state10,
                           'Histogram of differences of state 10 with diffe


# get the 20th state
gru_state20 = tf.transpose(gru_outputs, [1, 0, 2])[20]
show_hist_of_hidden_values(session, gru_state20,
                           'Histogram of differences of state 20 with diffe

# get the 20th state
gru_state30 = tf.transpose(gru_outputs, [1, 0, 2])[30]
show_hist_of_hidden_values(session, gru_state30,
                           'Histogram of differences of state 30 with diffe

```







## Backpropagation: vanishing gradients and exploding gradients

In the experiment, you will observe vanishing gradients and exploding gradients

```

In [372]: # Calculate gradient with respect to the initial state

# the gradient with respect to state 30 is [1, 1, ..., 1]. Propagate the gr
rnn_loss30 = tf.reduce_sum(rnn_state30)
rnn_gradh = tf.gradients([rnn_loss30], [initial_state])[0]


scale = 1.0
wt_h = (np.random.rand(hidden_size, hidden_size) - 0.5) * scale
wt_x = (np.random.rand(input_size, hidden_size) - 0.5) * scale
bias = (np.random.rand(hidden_size) - 0.5) * scale
set_rnn_params(rnn_cell, session, wt_h, wt_x, bias)

#####
# 5. Observe vanishing gradients (3 points)
# Try a different settings of the parameters and see if you still get vanis
#####

# show the norms of gradients. Most of them are zero.
np_rnn_gradh = session.run(rnn_gradh)
rnn_grad_norm = np.linalg.norm(np_rnn_gradh, axis=1)

n_bins = 20
_ = plt.hist(rnn_grad_norm, bins=n_bins)

```




```
In [376]: # Can you set GRU parameters such that the gradient does not vanish?

# the gradient with respect to state 30 is [1, 1, ..., 1]. Propagate the gr
gru_loss30 = tf.reduce_sum(gru_state30)
gru_gradh = tf.gradients([gru_loss30], [initial_state])[0]

#####
# 6. GRU parameters that don't have vanishing gradients (3 points)
# Set GRU parameters so that the gradient of a later state with respect to
# initial state is not near zero.
#####

# show norms of gradients
np_gru_gradh = session.run(gru_gradh)
gru_grad_norm = np.linalg.norm(np_gru_gradh, axis=1)

n_bins = 20
_ = plt.hist(gru_grad_norm, bins=n_bins)
```



In [ ]:

